

 Z.ai: GLM 5.2 (free)

z-ai/glm-5.2:free



 Model weights 

 Compare

 Playground

 Try this model

GLM 5.2 is a large-scale reasoning model from Z.ai. It supports text input and output with a 1M-token context window, and is suited for long-horizon agent workflows, project-level software engineering, and complex multi-step automation....

Show more 

MODALITIES

T → T

PRICE

Free

CONTEXT

256K

RELEASED

Jun 17, 2026

Providers


Different companies host the same model. OpenRouter routes your request to one of them based on the routing mode you pick — Balanced (price + speed), Nitro (fastest), or Exacto (highest tool-calling accuracy).

Free  		Latency / throughput				P50 
Provider 		 Input /M	 Output /M	 Latency	 Throughput	 Uptime
 Decart 		Free	Free	3.65s	99 tps	99.85%


Performance

Throughput is how fast the model writes (tokens per second — higher is better). Latency is total round-trip time (lower is better). TTFT is time-to-first-token — how long before you see anything appear (lower is better).


Throughput 99 tok/s P50, best across providers	Latency  3.65 s P50, best provider
All locations 	1 week 



P99	Avg 797 tok/s
P95	Avg 628 tok/s
P90	Avg 473 tok/s
P75	Avg 283 tok/s
P50	Avg 149 tok/s



P50	Avg 3.25 s
P75	Avg 5.71 s
P90	Avg 11.60 s
P95	Avg 17.32 s
P99	Avg 35.00 s




P50	Avg 8.91 s
P75	Avg 15.73 s
P90	Avg 32.84 s
P95	Avg 64.28 s
P99	Avg 120.49 s


AutoExacto Benchmarks

	GPQA Diamond	IAU-Bench
Firewor...	86.2%	--
lonstre...	85.2%	--
Wafer F...	92.2%	71.7%
Firewor...	85.4%	76.7%
Deba	88.5%	71.1%

Tool Call Error Rate




Structured Output Error Rate



Firewor...	86.1%	74.0%
Wafer	83.2%	73.8%
Makora	81.0%	75.3%

+29 more providers

Cache Hit Rate



Uptime is the percentage of the past 3 days that at least one provider was responding to requests.

Availability is the percentage of time that inference was successfully served. OpenRouter continuously monitors and uses the next-best provider when one returns an error.

Uptime (3d)

100.00%

Availability (3d)


96.34%

Availability over the last 3 days Aug 31, 6 PM - Sep 3, 6 PM

Availability 96.34%



Availability over the last 24 hours



When an error occurs in an upstream provider, we can recover by routing to another healthy provider, if your request filters allow it. You can access per-provider uptime data programmatically through the [Endpoints API](#). [Learn more](#) about our load balancing and customization options.

Benchmarks

Scores on standardized evaluations. Higher percentages are better — and rank percentile shows where this model lands among all models on OpenRouter.

 Artificial Analysis

 Design Arena

GLM-5.2 (max)

52.6

Artificial Analysis
Intelligence Index

Better than 79% of models
compared

68.8

Artificial Analysis
Coding Index

Better than 77% of models
compared

45.7

Artificial Analysis
Agentic Index

Better than 80% of models
compared

GPQA Diamond

Graduate-level scientific reasoning

89.5%

HLE

Humanity's Last Exam

41.1%

IFBench

Instruction-following benchmark

73.3%

τ^2 -Bench Telecom

Conversational AI agents in dual-control scenarios

99.1%

AA-LCR

Long context reasoning evaluation

76.7%

GDPval-AA

Economically valuable tasks

49.9%

CritPt

Research-level physics reasoning

20.9%

SciCode

Python programming for scientific computing

50.5%

Terminal-Bench Hard

Agentic coding & terminal use

50.8%

Knowledge

AA-Omniscience Accuracy

Proportion of correctly answered questions

24.3%

AA-Omniscience Non-Hallucination Rate

Rate of avoiding hallucination among non-correct responses

73.7%

Metrics sourced from [Artificial Analysis](#)

Apps

Public apps that send the most traffic to this model. Good signal for what real production workloads look like — and a hint at which use cases this model is best suited for.

- | | | |
|----|---|--------------|
| 1. |  <u>Janitor AI</u>
Janitor AI is a chatbot platform where users create and chat with custo... | 4.49B tokens |
| 2. |  <u>Kilo Code</u>
Kilo Code is an open-source AI coding agent that works across VS Co... | 3.05B tokens |
| 7. |  <u>Hermes Agent</u> | 2.07B tokens |


4.

 **SillyTavern**
SillyTavern is the LLM frontend for power users, a chat interface that co...

1.41B tokens
5.

 **Bary**
Organize and enhance your roleplay creations.

1.11B tokens




Activity

Token volume and request traffic to this model over time.

Tokens

3.8B





Quick Start

Drop-in code to call this model. OpenRouter's API is OpenAI-compatible — most SDKs work by just swapping the base URL. The only thing that changes between models is the model slug below.

1 Get your API key

Create an API key from your OpenRouter dashboard and set it as an environment variable:

Create API Key

```
export OPENROUTER_API_KEY=sk-or-v1- ...
```

[Copy](#)

2 Make your first request

Use `z-ai/glm-5.2:free` with the OpenRouter API:

internal reasoning before the final answer. When continuing a conversation, preserve the complete `reasoning_details` when passing messages back to the model so it can continue reasoning from where it left off. [Learn more about reasoning tokens](#).

In the examples below, the [OpenRouter-specific headers](#) are optional. Setting them allows your app to appear on the OpenRouter leaderboards.

[TypeScript SDK](#) [Python](#) [TypeScript \(fetch\)](#) [cURL](#) [Python \(OpenAI\)](#) [TypeScript \(OpenAI\)](#)

```
import { OpenRouter } from "@openrouter/sdk";

const openrouter = new OpenRouter({
  apiKey: "<OPENROUTER_API_KEY>"
});

// Stream the response to get reasoning tokens in usage
const stream = await openrouter.chat.send({
  chatRequest: {
    model: "z-ai/glm-5.2:free",
    messages: [
      {
        role: "user",
```

[Copy](#)

Using third-party SDKs

For information about using third-party SDKs and frameworks with OpenRouter, please see our [frameworks documentation](#).

3 Enable streaming

Add `"stream": true` to your request body to receive responses as server-sent events:

```
curl -N https://openrouter.ai/api/v1/chat/completions \
-H "Content-Type: application/json" \
```

[Copy](#)

```
{
  "model": "z-ai/glm-5.2:free",
  "stream": true,
  "messages": [
    { "role": "user", "content": "Hello" }
  ]
}
```

Endpoint

Sends a request for a model response for the given chat conversation. Supports both streaming and non-streaming modes.

POST https://openrouter.ai/api/v1/chat/completions

Authorization	Bearer \$OPENROUTER_API_KEY
Content-Type	application/json
HTTP-Referer	optional – your site URL, for rankings
X-Title	optional – your site name, for rankings
Model	z-ai/glm-5.2:free

Creates a streaming or non-streaming response using the OpenAI Responses API format. [Docs](#)

POST https://openrouter.ai/api/v1/responses

Authorization	Bearer \$OPENROUTER_API_KEY
Content-Type	application/json
HTTP-Referer	optional – your site URL, for rankings
X-Title	optional – your site name, for rankings

Creates a message using the Anthropic Messages API format. Supports text, images, PDFs, tools, and extended thinking. [Docs](#)

POST https://openrouter.ai/api/v1/messages

Authorization	Bearer \$OPENROUTER_API_KEY
Content-Type	application/json
HTTP-Referer	optional – your site URL, for rankings
X-Title	optional – your site name, for rankings
Model	z-ai/glm-5.2:free

Parameters

Name	Type	Default	Description
reasoning	map	—	Controls reasoning behavior for models that support thinking tokens, inc
temperature	float	1	This setting influences the variety in the model's responses.
top_p	float	0.95	This setting limits the model's choices to a percentage of likely tokens: o
top_k	integer	0	This limits the model's choice of tokens at each step, making it choose fr
min_p	float	0	Represents the minimum probability for a token to be considered, relativ
frequency_penalty	float	0	This setting aims to control the repetition of tokens based on how often
presence_penalty	float	0	Adjusts how often the model repeats specific tokens already used in the

stop	array	—	Stop generation immediately if the model encounter any token specified
seed	integer	—	If specified, the inferencing will sample deterministically, such that repeated
max_tokens	integer	—	This sets the upper limit for the number of tokens the model can generate
response_format	map	—	Forces the model to produce specific output format.
tools	array	—	Tool calling parameter, following OpenAI's tool calling request shape.
tool_choice	string or object	—	Controls which (if any) tool is called by the model.

Frequently asked questions

What is GLM 5.2 (free)?

Is GLM 5.2 (free) free?

What is the context length of GLM 5.2 (free)?

Does GLM 5.2 (free) support tool calling and structured outputs?

What other text models does Z.ai have?

When was GLM 5.2 (free) released?



Personal ▾

AI Model Rankings Ranking

More models from Z.ai

GLM Flash Latest

This model always redirects to the latest model in the GLM Flash family.

Text 1.0M context

GLM 5.3 Flash

GLM-5.3-Flash is a native multimodal model from Z.ai. It is suited for efficient ...

Text 1.3M context \$0.075 / 0.25

GLM 5.3 Flash

GLM-5.3-Flash is a native multimodal model from Z.ai. It is suited for efficient ...

Text 1.0M context \$0.15 / 0.50



© 2026 OpenRouter, Inc



Rankings

Benchmarks

Apps

Discover

Models

Collections

Providers

Pricing

Enterprise

Labs

Blog

Careers [Hiring](#)

Privacy

Terms of Service

Trust Center

Support

Works With OR

Data

Brand

API Reference

Developer Platform

Status

GitHub

LinkedIn

X

YouTube



Personal ▾



Personal ▾



Personal ▾



Personal ▾





Personal ▾







Personal ▾





