

 MiniMax: MiniMax M3 (free)

[minimax/minimax-m3:free](#)



 Model weights 

 Compare

 Playground

 Try this model

MiniMax-M3 is a multimodal foundation model from MiniMax. It supports text, image, and video inputs with text output, a 1M-token context window, and is suited for long-horizon agentic work, coding, and tool use. It is built on MiniMax Sparse ...

Show more 

MODALITIES

   → 

PRICE

Free

CONTEXT

1M

RELEASED

May 31, 2026

Providers

Different companies host the same model. OpenRouter routes your request to one of them based on the routing mode you pick — Balanced (price + speed), Nitro (fastest), or Exacto (highest tool-calling accuracy).

Free 

Latency / throughput

P50 ▾

Provider ↑↓

↑↓ Input /M

↑↓ Output /M

↑↓ Latency

↑↓ Throughput

↑↓ Uptime

GMICloud 

Free

Free

3.63s

25 tps

99.93%

⚡ Performance

Throughput is how fast the model writes (tokens per second — higher is better). Latency is total round-trip time (lower is better). TTFT is time-to-first-token — how long before you see anything appear (lower is better).

Throughput

25 tok/s

P50, best across providers

Latency ⓘ

3.63 s


P50, best provider

All locations ▾


1 week ▾



P99	Avg 179 tok/s
P95	Avg 129 tok/s
P90	Avg 103 tok/s
P75	Avg 66 tok/s
P50	Avg 36 tok/s



P50	Avg 2.85 s
P75	Avg 5.32 s
P90	Avg 10.44 s
P95	Avg 15.33 s
P99	Avg 32.70 s



P50	Avg 9.10 s
P75	Avg 17.25 s
P90	Avg 32.17 s
P95	Avg 47.83 s
P99	Avg 117.98 s

AutoExacto Benchmarks

GPQA DiamondTAU-Benchvgi_bench

89.1%76.5%--

88.6%74.1%--

88.6%74.1%

Tool Call Error Rate



Structured Output Error Rate

80.6% -- --

89.2% 71.1% --

+9 more providers

Cache Hit Rate



GMICloud Avg 85.93 %



Personal ▾

Uptime is the percentage of the past 3 days that at least one provider was responding to requests.

Availability is the percentage of time that inference was successfully served. OpenRouter continuously monitors and uses the next-best provider when one returns an error.

Uptime (3d)

100.00%

Availability (3d)

98.89%

Availability over the last 3 days Aug 31, 4 PM - Sep 3, 4 PM

Availability 98.89%




Mon

Tue

Wed

Now

Availability over the last 24 hours



When an error occurs in an upstream provider, we can recover by routing to another healthy provider, if your request filters allow it. You can access per-provider uptime data programmatically through the [Endpoints API](#). [Learn more](#) about our load balancing and customization options.

Benchmarks

Scores on standardized evaluations. Higher percentages are better — and rank percentile shows where this model lands among all models on OpenRouter.

 Artificial Analysis Design Arena

45.4

Artificial Analysis
Intelligence Index

Better than 67% of models
compared

58.6

Artificial Analysis
Coding Index

Better than 66% of models
compared

36.1

Artificial Analysis
Agentic Index

Better than 68% of models
compared

Reasoning

GPOA Diamond

Coding

SciCode

HLE

Humanity's Last Exam

39.0%

IFBench

Instruction-following benchmark

82.9%

τ^2 -Bench Telecom

Conversational AI agents in dual-control scenarios

88.9%

AA-LCR

Long context reasoning evaluation

80.3%

GDPval-AA

Economically valuable tasks

44.0%

CritPt

Research-level physics reasoning

3.7%

Terminal-Bench Hard

Agentic coding & terminal use

42.4%

Knowledge

Proportion of correctly answered questions

16.7%

AA-Omniscience Non-Hallucination Rate

Rate of avoiding hallucination among non-correct responses

81.6%

Metrics sourced from [Artificial Analysis](#)

Apps

Public apps that send the most traffic to this model. Good signal for what real production workloads look like — and a hint at which use cases this model is best suited for.

-  **[Hermes Agent](#)**
Hermes Agent is an open-source, self-improving AI agent by Nous Res... 588B tokens
-  **[Kilo Code](#)**
Kilo Code is an open-source AI coding agent that works across VS Co... 509B tokens
-  **[Claude Code](#)**
Claude Code is Anthropic's agentic coding tool that reads your entire ... 492B tokens
-  **[omp](#)** 17.5B


5.



DeepSeek Harness

new

156B tokens



Activity

Token volume and request traffic to this model over time.


Tokens

1T

750B

Prompt

273B



Quick Start

Drop-in code to call this model. OpenRouter's API is OpenAI-compatible — most SDKs work by just swapping the base URL. The only thing that changes between models is the model slug below.

1 Get your API key

Create an API key from your OpenRouter dashboard and set it as an environment variable:

Create API Key

```
export OPENROUTER_API_KEY=sk-or-v1- ...
```

[Copy](#)

2 Make your first request

Use `minimax/minimax-m3:free` with the OpenRouter API:

OpenRouter supports reasoning-enabled models that can show their step-by-step thinking process. Use the `reasoning` parameter in your request to enable reasoning, and access the `reasoning_details` array in the response to see the model's

In the examples below, the OpenRouter-specific headers are optional. Setting them allows your app to appear on the OpenRouter leaderboards.

TypeScript SDK Python TypeScript (fetch) cURL Python (OpenAI) TypeScript (OpenAI)

```
import { OpenRouter } from "@openrouter/sdk";

const openrouter = new OpenRouter({
  apiKey: "<OPENROUTER_API_KEY>"
});

// Stream the response to get reasoning tokens in usage
const stream = await openrouter.chat.send({
  chatRequest: {
    model: "minimax/minimax-m3:free",
    messages: [
      {
        role: "user",
```

[Copy](#)

Using third-party SDKs

For information about using third-party SDKs and frameworks with OpenRouter, please see our [frameworks documentation](#).

3 Enable streaming

Add `"stream": true` to your request body to receive responses as server-sent events:

```
curl -N https://openrouter.ai/api/v1/chat/completions \
  -H "Content-Type: application/json" \
  -H "Authorization: Bearer $OPENROUTER_API_KEY" \
  -d '{
    "model": "minimax/minimax-m3:free".
```

[Copy](#)

```
{ "role": "user", "content": "Hello" }  
]  
'
```

Endpoint

Sends a request for a model response for the given chat conversation. Supports both streaming and non-streaming modes.

POST https://openrouter.ai/api/v1/chat/completions

Authorization	Bearer \$OPENROUTER_API_KEY
Content-Type	application/json
HTTP-Referer	optional – your site URL, for rankings
X-Title	optional – your site name, for rankings
Model	minimax/minimax-m3:free

Creates a streaming or non-streaming response using the OpenAI Responses API format. [Docs](#)

POST https://openrouter.ai/api/v1/responses

Authorization	Bearer \$OPENROUTER_API_KEY
Content-Type	application/json
HTTP-Referer	optional – your site URL, for rankings
X-Title	optional – your site name, for rankings
Model	minimax/minimax-m3:free

POST https://openrouter.ai/api/v1/messages

Authorization Bearer \$OPENROUTER_API_KEY

Content-Type application/json

HTTP-Referer optional – your site URL, for rankings

X-Title optional – your site name, for rankings

Model minimax/minimax-m3:free

Parameters

Name	Type	Default	Description
reasoning	map	—	Controls reasoning behavior for models that support thinking tokens, includ
max_tokens	integer	—	This sets the upper limit for the number of tokens the model can generate in
temperature	float	1	This setting influences the variety in the model's responses.
top_p	float	0.95	This setting limits the model's choices to a percentage of likely tokens: only 1
seed	integer	—	If specified, the inferencing will sample deterministically, such that repeated
response_format	map	—	Forces the model to produce specific output format.
tools	array	—	Tool calling parameter, following OpenAI's tool calling request shape.
tool_choice	string or object	—	Controls which (if any) tool is called by the model.

Frequently asked questions

What is MiniMax M3 (free)?

Is MiniMax M3 (free) free?

What is the context length of MiniMax M3 (free)?

Does MiniMax M3 (free) support tool calling and structured outputs?

What inputs and outputs does MiniMax M3 (free) support?

What other text models does MiniMax have?

When was MiniMax M3 (free) released?

Explore more models

AI Models with Vision: Multimodal LLMs for Image Understanding Collection

AI Model Rankings Ranking

More models from MiniMax

H3 Max

MiniMax H3 Max is a video-generation model from MiniMax, jointly released with ...

Video from \$0.05/second

H3

MiniMax H3 is a lightweight, open-weights video generation model from MiniMax. It i...

Video from \$0.13/second

Speech 2.8 HD

MiniMax Speech 2.8 HD is a text-to-s model from MiniMax. It is suited for ...

Speech \$100/M characters



Personal ▾



© 2026 OpenRouter, Inc

Product

[Chat](#)

[Rankings](#)

[Benchmarks](#)

[Apps](#)

[Discover](#)

[Models](#)

[Collections](#)

[Providers](#)

[Pricing](#)

[Enterprise](#)

[Labs](#)

Company

[About](#)

[Blog](#)

[Careers](#) [Hiring](#)

[Privacy](#)

[Terms of Service](#)

[Trust Center](#)

[Support](#)

[Works With OR](#)

[Data](#)

[Brand](#)

Developer

[Documentation](#)

[API Reference](#)

[Developer Platform](#)

[Status](#)

Connect

[Discord](#)

[GitHub](#)

[LinkedIn](#)

[X](#)

[YouTube](#)



Personal ▾



Personal ▾



Personal ▾



Personal ▾



Personal ▾



Personal ▾



Personal ▾



Personal ▾



Personal ▾



Personal ▾



Personal ▾



Personal ▾